

# pdf2ipynb

July 2, 2021

## 1 Objective

As a good example of extracting information from PDF files, we set out to convert a Jupyter notebook in PDF form to its original `.ipynb` form.

I don't know if this is hard. Let's get started.

```
[1]: from pathlib import Path
     from PyPDF2 import PdfFileReader
```

N.B.

- The `n_pages` assignment should be kept in the `with` context; otherwise, one'd get a `ValueError: seek of closed file`
  - same as `pdf.getPage()`

```
[2]: path_pdf = Path("pdf2ipynb.pdf")
     with open(path_pdf, "rb") as f:
         pdf = PdfFileReader(f)
         n_pages = pdf.getNumPages()
         page00 = pdf.getPage(0)
         print(page00)
pdf
```

```
{'/Resources': IndirectObject(21, 0), '/Type': '/Page', '/Parent':
IndirectObject(69, 0), '/Contents': [IndirectObject(20, 0)], '/MediaBox': [0, 0,
612, 792]}
```

```
[2]: <PyPDF2.pdf.PdfFileReader at 0x7f0f180c9290>
```

```
[3]: n_pages
```

```
[3]: 8
```

```
[4]: info = pdf.getDocumentInfo()
     info
```

```
[4]: {'/Creator': 'LaTeX with hyperref',
      '/Producer': 'xdvipdfmx (20210318)',
      '/CreationDate': "D:20210702211923+07'00'"}

```

```
[5]: [ s for s in dir(pdf) if not s.startswith("_")]
```

```
[5]: ['cacheGetIndirectObject',  
      'cacheIndirectObject',  
      'decrypt',  
      'documentInfo',  
      'flattenedPages',  
      'getDestinationPageNumber',  
      'getDocumentInfo',  
      'getFields',  
      'getFormTextFields',  
      'getIsEncrypted',  
      'getNamedDestinations',  
      'getNumPages',  
      'getObject',  
      'getOutlines',  
      'getPage',  
      'getPageLayout',  
      'getPageMode',  
      'getPageNumber',  
      'getXmpMetadata',  
      'isEncrypted',  
      'namedDestinations',  
      'numPages',  
      'outlines',  
      'pageLayout',  
      'pageMode',  
      'pages',  
      'read',  
      'readNextEndLine',  
      'readObjectHeader',  
      'resolvedObjects',  
      'stream',  
      'strict',  
      'trailer',  
      'xmpMetadata',  
      'xref',  
      'xrefIndex',  
      'xref_objStm']
```

The above few cells came from the real python's tutorial, but only after reading the first few paragraphs of it did I find out that to extract the content of a PDF file, people seems to not recommend `pypdf2`; instead, people suggest using `pdfminer` (or `pdfminer.six`).

I chose to install `pip install pdfminer.six` because it seems to be a fork of `pdfminer` that is being constantly maintained, whereas `pdfminer` itself seems to be free of maintainance.

```
[6]: import pdfminer
      dir(pdfminer)
```

```
[6]: ['__builtins__',
      '__cached__',
      '__doc__',
      '__file__',
      '__loader__',
      '__name__',
      '__package__',
      '__path__',
      '__spec__',
      '__version__',
      'sys',
      'warnings']
```

```
[7]: from pdfminer.high_level import extract_text
      text = extract_text(path_pdf)
      text
```

```
[7]: 'pdf2ipynb\n\nJuly 2, 2021\n\n1 Objective\n\nAs a good example of extracting
information from PDF files, we set out to convert a Jupyter\nnotebook in PDF
form to its original .ipynb form.\n\n• The n_pages assignment should be kept in
the with context; otherwise, one'd get a\n\n[19]: path_pdf =
Path("JupyterNotebook-LinearRegression-MultipleInput.pdf")\n\nI don't know if
this is hard. Let's get started.\n\n[1]: from pathlib import Path\n\nfrom PyPDF2
import PdfFileReader\n\nN.B.\n\nValueError: seek of closed file\n\n- same as
pdf.getPage()\n\nwith open(path_pdf, "rb") as f:\nnpdf =
PdfFileReader(f)\nn_pages = pdf.getNumPages()\npage00 =
pdf.getPage(0)\nprint(page00)\n\nnpdf\n\n{\n  '/Resources\n': IndirectObject(19, 0),
  '/Type\n': \n'/Page\n', \n'/Parent\n':\nIndirectObject(54, 0), \n'/Contents\n':
[IndirectObject(18, 0)], \n'/MediaBox\n': [0, 0,\n612, 792]}\n\n[19]:
<PyPDF2.pdf.PdfFileReader at 0x7fecc84ab2d0>\n\n[17]: n_pages\n\n[17]: 5\n\n[6]:
info = pdf.getDocumentInfo()\n\ninfo\n\n[6]: {\n  '/Creator\n': \n'LaTeX with
hyperref package\n',\n\n\n'/Producer\n': \n'XeTeX 0.99998\n',\n\n\n'/CreationDate\n':
"D:20210625090544+07'\n00'\n"}\n\n1\n\n\nx0c[4]: [ s for s in dir(pdf) if not
s.startswith("_")]\n\n[4]: [\n'cacheGetIndirectObject\n',\n\n\n'\n'cacheIndirectObject
\n',\n\n\n'\n'decrypt\n',\n\n\n'\n'documentInfo\n',\n\n\n'\n'flattenedPages\n',\n\n\n'\n'getDestinationPageN
umber\n',\n\n\n'\n'getDocumentInfo\n',\n\n\n'\n'getFields\n',\n\n\n'\n'getFormTextFields\n',\n\n\n'\n'getIsE
ncrypted\n',\n\n\n'\n'getNamedDestinations\n',\n\n\n'\n'getNumPages\n',\n\n\n'\n'getObject\n',\n\n\n'\n'getO
utlines\n',\n\n\n'\n'getPage\n',\n\n\n'\n'getPageLayout\n',\n\n\n'\n'getPageMode\n',\n\n\n'\n'getPageNumber\n
',\n\n\n'\n'getXmpMetadata\n',\n\n\n'\n'isEncrypted\n',\n\n\n'\n'namedDestinations\n',\n\n\n'\n'numPages\n',
\n\n\n'\n'outlines\n',\n\n\n'\n'pageLayout\n',\n\n\n'\n'pageMode\n',\n\n\n'\n'pages\n',\n\n\n'\n'read\n',\n\n\n'\n'readNe
xtEndLine\n',\n\n\n'\n'readObjectHeader\n',\n\n\n'\n'resolvedObjects\n',\n\n\n'\n'stream\n',\n\n\n'\n'strict
\n',\n\n\n'\n'trailer\n',\n\n\n'\n'xmpMetadata\n',\n\n\n'\n'xref\n',\n\n\n'\n'xrefIndex\n',\n\n\n'\n'xref_objStm\n']
\n\nThe above few cells came from the real python's tutorial, but only after
reading the first few\nparagraphs of it did I find out that to extract the
```

content of a PDF file, people seems to not recommend pypdf2; instead, people suggest using pdfminer (or pdfminer.six). I chose to install pip install pdfminer.six because it seems to be a fork of pdfminer that is being constantly maintained, whereas pdfminer itself seems to be free of maintainance.

```

n2\n\nc[24]: import pdfminer\n\ndir(pdfminer)\n\n[24]: [\n'__b
uiltins__',\n\n'__cached__',\n\n'__doc__',\n\n'__file__',\n\n'__loader__',\n\n
'__name__',\n\n'__package__',\n\n'__path__',\n\n'__spec__',\n\n'__version__',\n\n
'sys',\n\n'warnings']\n\n[26]: from pdfminer.high_level import
extract_text\n\ntext = extract_text(path_pdf)\ntext\n\n[26]: \nLinearRegression-
MultipleInput-TensorBoard\n\n\nJune 25, 2021\n\n\n1 Import and\ncheck TensorFlow
version\n\n\n[14]: import numpy as np\n\n\nimport tensorflow as\ntf\n\nimport
matplotlib.pyplot as plt\n\nprint(tf.__version__)\n\n\n2.5.0\n\n\n2\nDownload and
check Boston Housing dataset\n\n\n[15]: from\ntensorflow.keras.datasets import
boston_housing\n\n\n(x_train, y_train), (x_test,\ny_test) =
boston_housing.load_data()\n\n\nndef normalize(x,y):\n\n\nnmean_y
=ny.mean(axis=0)\n\n\nstd_y = y.std(axis=0)\n\n\nnmean_x = x.mean(axis=0)\n\n\nstd_x
=nx.std(axis=0)\n\n\nnx_norm = (x - mean_x)/std_x\n\n\nny_norm = (y
-nmean_y)/std_y\n\n\n\nreturn x_norm, y_norm\n\n\nnx_train, y_train =
normalize(x_train,\ny_train)\n\n\nnx_test, y_test\n\n\nny_test)\n\n\nprint(\\n\\n\\n'Train:\\n\\n\\n',
x_train.shape,\ny_train.shape)\n\n\nprint(\\n\\n\\n'Test:\\n\\n\\n', x_test.shape,
y_test.shape)\n\n\n\n=normalize(x_test,\n\n\n\nTrain: (404, 13) (404,)\n\nTest:
(102, 13)\n\n(102,)\n\n\n1\n\n\n\n\nc3 Prepare the model:\n\n\nEquation: y = ax +
b\n\n- a,b:\n\nparameters that we need to find - x,y: observed data from the
reality\n\n\n[16]: a= tf.Variable(tf.random.uniform(shape=[13], minval=-1.,
maxval=1.))\n\n\nb =\ntf.Variable(tf.random.uniform(shape=[], minval=0.,
maxval=.5))\n\n\n\npredict(x):\n\n\n\nreturn tf.reduce_sum(a * x, 1) +
b\n\n\n\nmse(groundtruth,\n\n\nprediction):\n\n\n\nreturn
tf.reduce_mean(tf.square(groundtruth-
prediction))\n\n\n\n\ntest_the_model():\n\n\n\nprint(f\\n\\n\\n'Model:
a={np.around(a.numpy(),3)},\nb={b.numpy()}\n\\n\\n\\n')\n\n\nny_test_hat =
predict(x_test)\n\n\nerror = mse(y_test,\ny_test_hat)\n\n\nprint(\\n\\n\\n'Test MSE:\\n\\n\\n',\n\n
error.numpy())\n\n\n\nplt.figure(figsize=[5,5])\n\n\nplt.scatter(y_test,\ny_test_hat)\n\n
\nplt.title("House\\n\\n's Price
Prediction")\n\n\nplt.xlim([-6,6]);\n\n\nplt.ylim([-6,6])\n\n\nplt.xlabel("True
Price")\n\n\nplt.ylabel("Predicted\nPrice")\n\n\nplt.show()\n\n\n\n\ndef
train(learning_rate=1e-1, epochs=5):\n\n\n\nnglobal\n\n\n\nnglobal b\n\n\n\n\nwriter = tf.
summary.create_file_writer("./tmp/mylogs")\n\n\n\n\nwith\n\n\nwriter.as_default():\n\n\n\n
for i in range(epochs):\n\n\n\n\nwith\n\n\n\n3\n\n\n\n\nc7f.GradientTape(persistent=True)
as g:\n\n\n\n\nloss = mse(y_train,\n\n\npredict(x_train))\n\n\n\n\nprint(f"Epoch {i+1},
Loss: ", loss.numpy())\n\n\n\nngrad_a =\n\n\nng.gradient(loss, a)\n\n\nngrad_b =
g.gradient(loss, b)\n\n\n\n\nna.assign_sub(learning_rate\n*
grad_a)\n\n\n\nnb.assign_sub(learning_rate
*\nngrad_b)\n\n\n\n\n2\n\n\n\n\n\nc7f.summary.scalar("train-loss",
loss,\n\n\nstep=i)\n\n\n\nntf.summary.scalar("train-loss-2",
2*loss,\n\n\nstep=i)\n\n\n\n\nwriter.flush()\n\n\n\n\ntest_the_model()\n\n\n\n\nModel: a=[-0.933
-0.083 0.089\n-0.365 -0.718 -0.586
-0.367\n\n\n\n0.317\n\n\n\n0.525\n\n\n\n0.804\n\n\n\n-0.548 -0.492

```

```

-0.608],\nb=0.10259240865707397\\n\\nTest MSE: 2.6259625\\n\\n4 Train the
model\\n\\n[ ]:\ntrain(1e-1, epochs=50)\\n\\n3\\n\\n\\n\\x0c5 Test the model after
training\\n\\n\\n[20]:\ntest_the_model()\\n\\nModel: a=[-0.118 0.14 -0.012
0.111\n-0.265\\n\\n0.263\\n\\n0.02\\n\\n-0.381\\n\\n0.261 -0.15\\n\\n-0.215
0.088 -0.436],\nb=6.539222940915579e-09\\n\\nTest MSE: 0.24517176\\n\\n6 View
TensorBoard\\n\\nNote:\nYou should select this cell and call to Interrupt the
Kernel after finishing\nchecking Tensor-\\nBoard\\n\\n[ ]: !tensorboard --logdir
./tmp/mylogs\\n\\n4\\n\\n\\x0c7\nHomework\\n\\n1. Apply denormalization for the
testing function\\n2. Apply\nStochastic Gradient Descent for training:\\n\\n\\n•
1-batch\\n\\n• mini-batch\\n\\n3. Try\nother loss functions:\\n\\n\\n• Root Mean
Squared Error,\\n\\n• Mean Absolute Error,\\n\\n•\nCombined
RMSE+MAE+MSE\\n\\n5\\n\\n\\x0c'\n[27]: print(text)\n\nJune 25,
2021\\n\\nLinearRegression-MultipleInput-TensorBoard\\n\\n1 Import and check
TensorFlow version\\n\\n[14]: import numpy as np\n\nimport tensorflow as
tf\nimport matplotlib.pyplot as plt\nprint(tf.__version__)\n\n2.5.0\\n\\n2
Download and check Boston Housing dataset\\n\\n[15]: from
tensorflow.keras.datasets import boston_housing\n\n(x_train, y_train), (x_test,
y_test) = boston_housing.load_data()\n\nndef normalize(x,y):\n\nmean_y =
y.mean(axis=0)\nstd_y = y.std(axis=0)\nmean_x =
x.mean(axis=0)\nstd_x = x.std(axis=0)\n\nx_norm = (x -
mean_x)/std_x\ny_norm = (y - mean_y)/std_y\n\nreturn x_norm, y_norm\n\nx_train,
y_train = normalize(x_train, y_train)\nx_test,
y_test\nny_test)\nprint('Train:', x_train.shape,
y_train.shape)\nprint('Test:', x_test.shape, y_test.shape)\n\n=
normalize(x_test,\n\nTrain: (404, 13) (404,)\nTest: (102, 13) (102,)\n\n1\n\n3
Prepare the model:\\n\\nEquation:  $y = ax + b$ \n- a,b: parameters that we need to
find - x,y: observed data from the reality\\n\\n[16]: a =
tf.Variable(tf.random.uniform(shape=[13], minval=-1., maxval=1.))\n\nb =
tf.Variable(tf.random.uniform(shape=[], minval=0., maxval=.5))\n\nndef
predict(x):\n\nreturn tf.reduce_sum(a * x, 1) + b\n\nndef mse(groundtruth,
prediction):\n\nreturn tf.reduce_mean(tf.square(groundtruth-prediction))\n\nndef
test_the_model():\n\nprint(f'Model: a={np.around(a.numpy(),3)},
b={b.numpy()}\nny_test_hat = predict(x_test)\nerror = mse(y_test,
y_test_hat)\nprint('Test MSE:',
error.numpy())\n\nplt.figure(figsize=[5,5])\nplt.scatter(y_test,
y_test_hat)\nplt.title("House's Price Prediction")\nplt.xlim([-6,6]);
plt.ylim([-6,6])\n\n5\\n\\n\\x0cwriter =
tf.summary.create_file_writer("./tmp/mylogs")\n\nplt.xlabel("True
Price")\nplt.ylabel("Predicted Price")\nplt.show()\n\nndef
train(learning_rate=1e-1, epochs=5):\n\nglobal a\nglobal b\n\nwith
writer.as_default():\n\nfor i in range(epochs):\n\nwith
tf.GradientTape(persistent=True) as g:\n\nloss = mse(y_train,
predict(x_train))\n\nprint(f'Epoch {i+1}, Loss: ", loss.numpy())\n\ngrad_a =
g.gradient(loss, a)\ngrad_b = g.gradient(loss, b)\n\nna.assign_sub(learning_rate
* grad_a)\nb.assign_sub(learning_rate * grad_b)\n\nntf.summary.scalar("train-
loss", loss, step=i)\nntf.summary.scalar("train-loss-2", 2*loss,
step=i)\n\nwriter.flush()\n\nntest_the_model()\n\nModel: a=[-0.933 -0.083 0.089

```

```
-0.365 -0.718 -0.586 -0.367\n\n2\n\n0.317\n\n0.525\n\n0.804\n\n-0.548 -0.492
-0.608], b=0.10259240865707397\n\nTest MSE: 2.6259625\n\n4 Train the
model\n\n6\n\nx0c3\n\n4\n\n5\n\n[ ]: train(1e-1, epochs=50)\n\n5 Test the model
after training\n\n[20]: test_the_model()\n\nModel: a=[-0.118 0.14 -0.012 0.111
-0.265\n\n0.263\n\n0.02\n\n-0.381\n\n0.261 -0.15\n\n-0.215 0.088 -0.436],
b=6.539222940915579e-09\n\nTest MSE: 0.24517176\n\n6 View TensorBoard\n\nNote:
You should select this cell and call to Interrupt the Kernel after\nfinishing
checking Tensor-\nBoard\n\n[ ]: !tensorboard --logdir ./tmp/mylogs\n\n1. Apply
denormalization for the testing function\n2. Apply Stochastic Gradient Descent
for training:\n\n7 Homework\n\n• 1-batch\n• mini-batch\n\n3. Try other loss
functions:\n\n• Root Mean Squared Error,\n• Mean Absolute Error,\n• Combined
RMSE+MAE+MSE\n\n7\n\nx0c1.1 Ref.\n\n[ ]:\n\n[ ]:\n\n•
https://realpython.com/pdf-python/\n• https://pdfminersix.readthedocs.io/en/latest/tutorial/highlevel.html\n\n8\n\nx0c'
```

```
[8]: print(text)
```

pdf2ipynb

July 2, 2021

## 1 Objective

As a good example of extracting information from PDF files, we set out to convert a Jupyter notebook in PDF form to its original .ipynb form.

- The `n_pages` assignment should be kept in the `with` context; otherwise, one'd get a

```
[19]: path_pdf = Path("JupyterNotebook-LinearRegression-MultipleInput.pdf")
```

I don't know if this is hard. Let's get started.

```
[1]: from pathlib import Path
```

```
from PyPDF2 import PdfFileReader
```

N.B.

```
ValueError: seek of closed file
```

```
- same as pdf.getPage()
```

```
with open(path_pdf, "rb") as f:
    pdf = PdfFileReader(f)
    n_pages = pdf.getNumPages()
    page00 = pdf.getPage(0)
```

```

print(page00)

pdf

{'/Resources': IndirectObject(19, 0), '/Type': '/Page', '/Parent':
IndirectObject(54, 0), '/Contents': [IndirectObject(18, 0)], '/MediaBox': [0, 0,
612, 792]}

[19]: <PyPDF2.pdf.PdfFileReader at 0x7fecc84ab2d0>

[17]: n_pages

[17]: 5

[6]: info = pdf.getDocumentInfo()

info

[6]: {'/Creator': 'LaTeX with hyperref package',

'/Producer': 'XeTeX 0.99998',
'/CreationDate': "D:20210625090544+07'00'"}

1

[4]: [ s for s in dir(pdf) if not s.startswith("_")]

[4]: ['cacheGetIndirectObject',

'cacheIndirectObject',
'decrypt',
'documentInfo',
'flattenedPages',
'getDestinationPageNumber',
'getDocumentInfo',
'getFields',
'getFormTextFields',
'getIsEncrypted',
'getNamedDestinations',
'getNumPages',
'getObject',
'getOutlines',
'getPage',
'getPageLayout',
'getPageMode',
'getPageNumber',
'getXmpMetadata',
'isEncrypted',

```

```

'namedDestinations',
'numPages',
'outlines',
'pageLayout',
'pageMode',
'pages',
'read',
'readNextEndLine',
'readObjectHeader',
'resolvedObjects',
'stream',
'strict',
'trailer',
'xmpMetadata',
'xref',
'xrefIndex',
'xref_objStm']

```

The above few cells came from the real python's tutorial, but only after reading the first few paragraphs of it did I find out that to extract the content of a PDF file, people seems to not recommend pypdf2; instead, people suggest using pdfminer (or pdfminer.six).

I chose to install pip install pdfminer.six because it seems to be a fork of pdfminer that is being constantly maintained, whereas pdfminer itself seems to be free of maintainance.

2

```

[24]: import pdfminer

dir(pdfminer)

[24]: ['__builtins__',

'__cached__',
'__doc__',
'__file__',
'__loader__',
'__name__',
'__package__',
'__path__',
'__spec__',
'__version__',
'sys',
'warnings']

```



```
[26]: from pdfminer.high_level import extract_text
```

```
text = extract_text(path_pdf)
text
```

```
[26]: 'LinearRegression-MultipleInput-TensorBoard\n\nJune 25, 2021\n\n1 Import
and
check TensorFlow version\n\n[14]: import numpy as np\n\nimport tensorflow as
tf\nimport matplotlib.pyplot as plt\nprint(tf.__version__)\n\n2.5.0\n\n2
Download and check Boston Housing dataset\n\n[15]: from
tensorflow.keras.datasets import boston_housing\n\n(x_train, y_train), (x_test,
y_test) = boston_housing.load_data()\n\nndef normalize(x,y):\n\nnmean_y =
y.mean(axis=0)\n\nstd_y = y.std(axis=0)\n\nnmean_x = x.mean(axis=0)\n\nstd_x =
x.std(axis=0)\n\nnx_norm = (x - mean_x)/std_x\n\nny_norm = (y -
mean_y)/std_y\n\nnreturn x_norm, y_norm\n\nnx_train, y_train = normalize(x_train,
y_train)\n\nnx_test, y_test = normalize(x_test, y_test)\n\nprint('\nTrain:\n', x_train.shape,
y_train.shape)\n\nprint('\nTest:\n', x_test.shape, y_test.shape)\n\nn=
normalize(x_test,\n\nnTrain: (404, 13) (404,)\n\nTest: (102, 13)
(102,)\n\n1\n\n\nx0c3 Prepare the model:\n\nnEquation:  $y = ax + b$ \n- a,b:
parameters that we need to find - x,y: observed data from the reality\n\n[16]: a
= tf.Variable(tf.random.uniform(shape=[13], minval=-1., maxval=1.))\n\nnb =
tf.Variable(tf.random.uniform(shape=[], minval=0., maxval=.5))\n\nndef
predict(x):\n\nnreturn tf.reduce_sum(a * x, 1) + b\n\nndef mse(groundtruth,
prediction):\n\nnreturn tf.reduce_mean(tf.square(groundtruth-prediction))\n\nndef
test_the_model():\n\nnprint(f'\nModel: a={np.around(a.numpy(),3)},
b={b.numpy()}\n')\n\nny_test_hat = predict(x_test)\n\nnerror = mse(y_test,
y_test_hat)\n\nprint('\nTest MSE:\n',
error.numpy())\n\nnplt.figure(figsize=[5,5])\n\nnplt.scatter(y_test,
y_test_hat)\n\nnplt.title("House's Price Prediction")\n\nnplt.xlim([-6,6]);
plt.ylim([-6,6])\n\nnplt.xlabel("True Price")\n\nnplt.ylabel("Predicted
Price")\n\nnplt.show()\n\nndef train(learning_rate=1e-1, epochs=5):\n\nnglobal
a\n\nnglobal b\n\nnwriter = tf.summary.create_file_writer("./tmp/mylogs")\n\nnwith
writer.as_default():\n\nnfor i in range(epochs):\n\nnwith
```

3

```
tf.GradientTape(persistent=True) as g:\n\nnloss = mse(y_train,
predict(x_train))\n\nnprint(f"Epoch {i+1}, Loss: ", loss.numpy())\n\nngrad_a =
g.gradient(loss, a)\n\nngrad_b = g.gradient(loss, b)\n\nna.assign_sub(learning_rate
* grad_a)\n\nnb.assign_sub(learning_rate *
grad_b)\n\nn2\n\n\nx0ctf.summary.scalar("train-loss", loss,
step=i)\n\nntf.summary.scalar("train-loss-2", 2*loss,
step=i)\n\nnwriter.flush()\n\nntest_the_model()\n\nnModel: a=[-0.933 -0.083 0.089
-0.365 -0.718 -0.586 -0.367\n\nn0.317\n\nn0.525\n\nn0.804\n\nn-0.548 -0.492 -0.608],
b=0.10259240865707397\n\nnTest MSE: 2.6259625\n\nn4 Train the model\n\nn[ ]:
train(1e-1, epochs=50)\n\nn3\n\n\nx0c5 Test the model after training\n\nn[20]:
test_the_model()\n\nnModel: a=[-0.118 0.14 -0.012 0.111
```

```
-0.265\n\n0.263\n\n0.02\n\n-0.381\n\n0.261 -0.15\n\n-0.215 0.088 -0.436],
b=6.539222940915579e-09\n\nTest MSE: 0.24517176\n\n6 View TensorBoard\n\nNote:
You should select this cell and call to Interrupt the Kernel after finishing
checking Tensor-\nBoard\n\n[ ]: !tensorboard --logdir ./tmp/mylogs\n\n4\n\n\x0c7
Homework\n\n1. Apply denormalization for the testing function\n2. Apply
Stochastic Gradient Descent for training:\n\n• 1-batch\n• mini-batch\n\n3. Try
other loss functions:\n\n• Root Mean Squared Error,\n• Mean Absolute Error,\n•
Combined RMSE+MAE+MSE\n\n5\n\n\x0c'
```

```
[27]: print(text)
```

June 25, 2021

LinearRegression-MultipleInput-TensorBoard

1 Import and check TensorFlow version

```
[14]: import numpy as np
```

```
import tensorflow as tf
import matplotlib.pyplot as plt
print(tf.__version__)
```

2.5.0

2 Download and check Boston Housing dataset

```
[15]: from tensorflow.keras.datasets import boston_housing
```

```
(x_train, y_train), (x_test, y_test) = boston_housing.load_data()
```

```
def normalize(x,y):
```

```
    mean_y = y.mean(axis=0)
    std_y = y.std(axis=0)
```

```
    mean_x = x.mean(axis=0)
```

4

```
    std_x = x.std(axis=0)
```

```
    x_norm = (x - mean_x)/std_x
    y_norm = (y - mean_y)/std_y
```

```
    return x_norm, y_norm
```

```
x_train, y_train = normalize(x_train, y_train)
```

```

x_test, y_test
y_test)
print('Train:', x_train.shape, y_train.shape)
print('Test:', x_test.shape, y_test.shape)

= normalize(x_test,

Train: (404, 13) (404,)
Test: (102, 13) (102,)

1

3 Prepare the model:

Equation:  $y = ax + b$ 
- a,b: parameters that we need to find - x,y: observed data from the reality

[16]: a = tf.Variable(tf.random.uniform(shape=[13], minval=-1., maxval=1.))

b = tf.Variable(tf.random.uniform(shape=[], minval=0., maxval=.5))

def predict(x):

return tf.reduce_sum(a * x, 1) + b

def mse(groundtruth, prediction):

return tf.reduce_mean(tf.square(groundtruth-prediction))

def test_the_model():

print(f'Model: a={np.around(a.numpy(),3)}, b={b.numpy()}')
y_test_hat = predict(x_test)
error = mse(y_test, y_test_hat)
print('Test MSE:', error.numpy())

plt.figure(figsize=[5,5])
plt.scatter(y_test, y_test_hat)
plt.title("House's Price Prediction")
plt.xlim([-6,6]); plt.ylim([-6,6])

5

writer = tf.summary.create_file_writer("./tmp/mylogs")

plt.xlabel("True Price")
plt.ylabel("Predicted Price")
plt.show()

```

```

def train(learning_rate=1e-1, epochs=5):

    global a
    global b

    with writer.as_default():

        for i in range(epochs):

            with tf.GradientTape(persistent=True) as g:

                loss = mse(y_train, predict(x_train))

                print(f"Epoch {i+1}, Loss: ", loss.numpy())

                grad_a = g.gradient(loss, a)
                grad_b = g.gradient(loss, b)

                a.assign_sub(learning_rate * grad_a)
                b.assign_sub(learning_rate * grad_b)

            tf.summary.scalar("train-loss", loss, step=i)

            tf.summary.scalar("train-loss-2", 2*loss, step=i)
            writer.flush()

        test_the_model()

Model: a=[-0.933 -0.083 0.089 -0.365 -0.718 -0.586 -0.367

2

0.317

0.525

0.804

-0.548 -0.492 -0.608], b=0.10259240865707397

Test MSE: 2.6259625

4 Train the model

6

3

```

4

5

```
[ ]: train(1e-1, epochs=50)
```

5 Test the model after training

```
[20]: test_the_model()
```

```
Model: a=[-0.118 0.14 -0.012 0.111 -0.265
```

```
0.263
```

```
0.02
```

```
-0.381
```

```
0.261 -0.15
```

```
-0.215 0.088 -0.436], b=6.539222940915579e-09
```

```
Test MSE: 0.24517176
```

6 View TensorBoard

Note: You should select this cell and call to Interrupt the Kernel after finishing checking TensorBoard

```
[ ]: !tensorboard --logdir ./tmp/mylogs
```

1. Apply denormalization for the testing function
2. Apply Stochastic Gradient Descent for training:

7 Homework

- 1-batch
- mini-batch

3. Try other loss functions:

- Root Mean Squared Error,
- Mean Absolute Error,
- Combined RMSE+MAE+MSE

7

1.1 Ref.

[ ]:

[ ]:

- <https://realpython.com/pdf-python/>
- <https://pdfminersix.readthedocs.io/en/latest/tutorial/highlevel.html>

8

[ ]:

1.1 Ref.

- <https://realpython.com/pdf-python/>
- <https://pdfminersix.readthedocs.io/en/latest/tutorial/highlevel.html>

[ ]: